

LAB 2

```
#include <linux/version.h>
#include <linux/autocfg.h>
#include <linux/module.h>
#include <linux/moduleparam.h>
#include <linux/init.h>

#include <linux/sched.h>
#include <linux/kernel.h> /* printk() */
#include <linux/errno.h> /* error codes */
#include <linux/types.h> /* size_t */
#include <linux/vmalloc.h>
#include <linux/genhd.h>
#include <linux/blkdev.h>
#include <linux/wait.h>
#include <linux/file.h>

#include "spinlock.h"
#include "osprd.h"

/* The size of an OSPRD sector. */
#define SECTOR_SIZE 512

/* This flag is added to an OSPRD file's f_flags to indicate that the file
 * is locked. */
#define F_OSPRD_LOCKED 0x80000

/* eprintk() prints messages to the console.
 * (If working on a real Linux machine, change KERN_NOTICE to KERN_ALERT or
 * KERN_EMERG so that you are sure to see the messages. By default, the
 * kernel does not print all messages to the console. Levels like KERN_ALERT
 * and KERN_EMERG will make sure that you will see messages.) */
#define eprintk(format, ...) printk(KERN_NOTICE format, ## __VA_ARGS__)

MODULE_LICENSE("Dual BSD/GPL");
MODULE_DESCRIPTION("CS 111 RAM Disk");
// EXERCISE: Pass your names into the kernel as the module's authors.
MODULE_AUTHOR("JONATHAN WOONG");

#define OSPRD_MAJOR 222

/* This module parameter controls how big the disk will be.
 * You can specify module parameters when you load the module,
 * as an argument to insmod: "insmod osprd.ko nsectors=4096" */
static int nsectors = 32;
module_param(nsectors, int, 0);

/* The internal representation of our device. */
typedef struct osprd_info {
    uint8_t *data; /* The data array. Its size is
                     // (nsectors * SECTOR_SIZE) bytes.

    osp_spinlock_t mutex; /* Mutex for synchronizing access to
                           // this block device

    unsigned ticket_head; /* Currently running ticket for
                           // the device lock

    unsigned ticket_tail; /* Next available ticket for
                           // the device lock

    wait_queue_head_t blockq; /* Wait queue for tasks blocked on
                              // the device lock

    size_t n_readl; /* Number of read locks

    size_t n_writel; /* Number of write locks

    unsigned desync; /* number of interrupted processes

    int dead; /* whether this will cause a deadlock or not
    /* HINT: You may want to add additional fields to help
    in detecting deadlock. */

    // The following elements are used internally; you don't need
    // to understand them.
    struct request_queue *queue; /* The device request queue.
    spinlock_t qlock; /* Used internally for mutual
                       // exclusion in the 'queue'.
    struct gendisk *gd; /* The generic disk.
    } osprd_info_t;

#define NOSPRD 4
static osprd_info_t osprds[NOSPRD];

// Declare useful helper functions

/*
 * file2osprd(filp)
 * Given an open file, check whether that file corresponds to an OSP ramdisk.
 * If so, return a pointer to the ramdisk's osprd_info_t.
 * If not, return NULL.
 */
static osprd_info_t *file2osprd(struct file *filp);

/*
 * for_each_open_file(task, callback, user_data)
 * Given a task, call the function 'callback' once for each of 'task's open
 * files. 'callback' is called as 'callback(filp, user_data)'; 'filp' is
 * the open file, and 'user_data' is copied from for_each_open_file's third
 * argument.
 */
static void for_each_open_file(struct task_struct *task,
                              void (*callback)(struct file *filp,
                                                  osprd_info_t *user_data),
                              osprd_info_t *user_data);

/*
```

LAB 2

```
* osprd_process_request(d, req)
*   Called when the user reads or writes a sector.
*   Should perform the read or write, as appropriate.
*/
static void osprd_process_request(osprd_info_t *d, struct request *req)
{
    long int sector_offset;
    long int numbytes;
    //int err;
    if (!blk_fs_request(req)) {
        end_request(req, 0);
        return;
    }

    // EXERCISE: Perform the read or write request by copying data between
    // our data array and the request's buffer.
    // Hint: The 'struct request' argument tells you what kind of request
    // this is, and which sectors are being read or written.
    // Read about 'struct request' in <linux/blkdev.h>.
    // Consider the 'req->sector', 'req->current_nr_sectors', and
    // 'req->buffer' members, and the rq_data_dir() function.

    // Your code here.

    sector_offset = req->sector*SECTOR_SIZE;
    numbytes = req->current_nr_sectors*SECTOR_SIZE;

    if(rq_data_dir(req)==READ) {
        memcpy(req->buffer, d->data+sector_offset, numbytes);
    }
    else if(rq_data_dir(req)==WRITE) {
        memcpy(d->data+sector_offset, req->buffer, numbytes);
    }
    else {
        eprintk("Failure to READ/WRITE\n");
    }
    end_request(req, 1);
}

// This function is called when a /dev/osprdX file is opened.
// You aren't likely to need to change this.
static int osprd_open(struct inode *inode, struct file *filp)
{
    // Always set the O_SYNC flag. That way, we will get writes immediately
    // instead of waiting for them to get through write-back caches.
    filp->f_flags |= O_SYNC;
    return 0;
}

// This function is called when a /dev/osprdX file is finally closed.
// (If the file descriptor was dup2ed, this function is called only when the
// last copy is closed.)
```

```
int osprd_ioctl(struct inode *inode, struct file *filp,
                unsigned int cmd, unsigned long arg);

static int osprd_close_last(struct inode *inode, struct file *filp)
{
    int r;
    r = 0;
    if (filp) {
        osprd_info_t *d = file2osprd(filp);
        int filp_writable = filp->f_mode & FMODE_WRITE;

        // EXERCISE: If the user closes a ramdisk file that holds
        // a lock, release the lock. Also wake up blocked processes
        // as appropriate.

        // Your code here.

        // This line avoids compiler warnings; you may remove it.
        (void) filp_writable, (void) d;
        if (!(filp->f_flags & F_OSPRD_LOCKED))
        {
            r = -EINVAL;
        }

        // Otherwise, clear the lock from filp->f_flags, wake up
        // the wait queue, perform any additional accounting steps
        // you need, and return 0.
        else
        {
            osp_spin_lock(&(d->mutex));
            filp->f_flags &= ~F_OSPRD_LOCKED;
            d->n_writel = 0;
            d->n_readl = 0;
            d->dead = 0;

            if(waitqueue_active(&d->blockq) == 0) {
                d->ticket_tail += d->desync;
                d->desync = 0;
            }
            osp_spin_unlock(&(d->mutex));
            wake_up_all(&d->blockq);
            r = 0;
        }
    }

    return r;
}

void cause_deadlock(struct file *filp, osprd_info_t *d){
    if (file2osprd(filp) == d){
        d->dead++;
    }
}

int osprd_ioctl(struct inode *inode, struct file *filp,
```

LAB 2

```
    unsigned int cmd, unsigned long arg)
{
    osprd_info_t *d = file2osprd(filp); // device info
    int r = 0;
    unsigned local_ticket;
    // is file open for writing?
    int filp_writable = (filp->f_mode & FMODE_WRITE) != 0;

    // This line avoids compiler warnings; you may remove it.
    // (void) filp_writable, (void) d;

    // Set 'r' to the ioctl's return value: 0 on success, negative on error

    if (cmd == OSPRDIOACQUIRE) {

        // EXERCISE: Lock the ramdisk.g
        //
        // If *filp is open for writing (filp_writable), then attempt
        // to write-lock the ramdisk; otherwise attempt to read-lock
        // the ramdisk.
        //
        // This lock request must block using 'd->blockq' until:
        // 1) no other process holds a write lock;
        // 2) either the request is for a read lock, or no other process
        //    holds a read lock; and
        // 3) lock requests should be serviced in order, so no process
        //    that blocked earlier is still blocked waiting for the
        //    lock.
        //
        // If a process acquires a lock, mark this fact by setting
        // 'filp->f_flags |= F_OSPRD_LOCKED'. You also need to
        // keep track of how many read and write locks are held:
        // change the 'osprd_info_t' structure to do this.
        //
        // Also wake up processes waiting on 'd->blockq' as needed.
        //
        // If the lock request would cause a deadlock, return -EDEADLK.
        // If the lock request blocks and is awoken by a signal, then
        // return -ERESTARTSYS.
        // Otherwise, if we can grant the lock request, return 0.

        // 'd->ticket_head' and 'd->ticket_tail' should help you
        // service lock requests in order. These implement a ticket
        // order: 'ticket_tail' is the next ticket, and 'ticket_head'
        // is the ticket currently being served. You should set a local
        // variable to 'd->ticket_head' and increment 'd->ticket_head'.
        // Then, block at least until 'd->ticket_tail == local_ticket'.
        // (Some of these operations are in a critical section and must
        // be protected by a spinlock; which ones?)

        // Your code here (instead of the next two lines).

        osp_spin_lock(&(d->mutex));
        local_ticket = d->ticket_head;
        d->ticket_head++;

        osp_spin_unlock(&(d->mutex));

        for_each_open_file(current, cause_deadlock, d);

        if (d->dead > 1 && (filp->f_flags & F_OSPRD_LOCKED)) {
            return -EDEADLK;
        }

        if (wait_event_interruptible(d->blockq, d->n_writel == 0
            && (!filp_writable || d->n_readl == 0)
            && d->ticket_tail == local_ticket))
        {
            if (d->ticket_tail == local_ticket) {
                d->ticket_tail++;
            }
            else {
                d->desync++;
            }
            return -ERESTARTSYS;
        }
        osp_spin_lock(&(d->mutex));

        d->dead = 0;

        if (d->mutex.lock > 0) {
            r = 0;
        }
        filp->f_flags |= F_OSPRD_LOCKED;
        if (filp_writable) {
            d->n_writel++; d->ticket_tail++;
        }
        else {
            d->n_readl++;
        }

        osp_spin_unlock(&(d->mutex));

        if (!filp_writable) {
            d->ticket_tail++;
        }

        r = 0;
    } else if (cmd == OSPRDIOCTRYACQUIRE) {

        // EXERCISE: ATTEMPT to lock the ramdisk.
        //
        // This is just like OSPRDIOACQUIRE, except it should never
        // block. If OSPRDIOACQUIRE would block or return deadlock,
        // OSPRDIOCTRYACQUIRE should return -EBUSY.
        // Otherwise, if we can grant the lock request, return 0.

        local_ticket = d->ticket_head;
        if (filp->f_flags & F_OSPRD_LOCKED || d->n_writel != 0
            || (filp_writable && d->n_readl != 0))
```

LAB 2

```

    ll d->ticket_tail != local_ticket)
{
    r = -EBUSY;
}

else
{
    osp_spin_lock(&(d->mutex));
    d->ticket_head++;
    filp->f_flags |= F_OSPRD_LOCKED;
    if (filp_writable) {
        d->n_writel++;
    }
    else {
        d->n_readl++;
    }
    if(d->ticket_tail < d->ticket_head) {
        d->ticket_tail++;
    }
    osp_spin_unlock(&(d->mutex));
    r = 0;
    wake_up_all(&d->blockq);
}

} else if (cmd == OSPRDIOCRELEASE) {

    // EXERCISE: Unlock the ramdisk.
    //
    // If the file hasn't locked the ramdisk, return -EINVAL.
    if (!filp->f_flags & F_OSPRD_LOCKED)) {
        r = -EINVAL;
    }

    // Otherwise, clear the lock from filp->f_flags, wake up
    // the wait queue, perform any additional accounting steps
    // you need, and return 0.
    else
    {
        osp_spin_lock(&(d->mutex));
        filp->f_flags &= ~F_OSPRD_LOCKED;

        d->n_writel = 0;
        d->n_readl = 0;

        osp_spin_unlock(&(d->mutex));

        wake_up_all(&d->blockq);

        r = 0;
    }
}
else {
    r = -ENOTTY;
}

```

```

    return r;
}

// Initialize internal fields for an osprd_info_t.

static void osprd_setup(osprd_info_t *d)
{
    /* Initialize the wait queue. */
    init_waitqueue_head(&d->blockq);
    osp_spin_lock_init(&d->mutex);
    d->ticket_head = d->ticket_tail = 0;
    d->n_readl = 0;
    d->n_writel = 0;
    d->dead = 0;
    d->desync = 0;
    /* Add code here if you add fields to osprd_info_t. */
}

/*****
/*   THERE IS NO NEED TO UNDERSTAND ANY CODE BELOW THIS LINE!   */
/*   */
*****/

// Process a list of requests for a osprd_info_t.
// Calls osprd_process_request for each element of the queue.

static void osprd_process_request_queue(request_queue_t *q)
{
    osprd_info_t *d = (osprd_info_t *) q->queuedata;
    struct request *req;

    while ((req = elv_next_request(q)) != NULL)
        osprd_process_request(d, req);
}

// Some particularly horrible stuff to get around some Linux issues:
// the Linux block device interface doesn't let a block device find out
// which file has been closed. We need this information.

static struct file_operations osprd_blk_fops;
static int (*blkdev_release)(struct inode *, struct file *);

static int _osprd_release(struct inode *inode, struct file *filp)
{
    if (file2osprd(filp))
        osprd_close_last(inode, filp);
    return (*blkdev_release)(inode, filp);
}

static int _osprd_open(struct inode *inode, struct file *filp)
{
    if (!osprd_blk_fops.open) {

```

LAB 2

```
    memcpy(&osprd_blk_fops, filp->f_op, sizeof(osprd_blk_fops));
    blkdev_release = osprd_blk_fops.release;
    osprd_blk_fops.release = _osprd_release;
}
filp->f_op = &osprd_blk_fops;
return osprd_open(inode, filp);
}
```

// The device operations structure.

```
static struct block_device_operations osprd_ops = {
    .owner = THIS_MODULE,
    .open = _osprd_open,
    // .release = osprd_release, // we must call our own release
    .ioctl = osprd_ioctl
};
```

// Given an open file, check whether that file corresponds to an OSP ramdisk.
// If so, return a pointer to the ramdisk's osprd_info_t.
// If not, return NULL.

```
static osprd_info_t *file2osprd(struct file *filp)
{
    if (filp) {
        struct inode *ino = filp->f_dentry->d_inode;
        if (ino->i_bdev
            && ino->i_bdev->bd_disk
            && ino->i_bdev->bd_disk->major == OSPRD_MAJOR
            && ino->i_bdev->bd_disk->fops == &osprd_ops)
            return (osprd_info_t *) ino->i_bdev->bd_disk->private_data;
    }
    return NULL;
}
```

// Call the function 'callback' with data 'user_data' for each of 'task's
// open files.

```
static void for_each_open_file(struct task_struct *task,
    void (*callback)(struct file *filp, osprd_info_t *user_data),
    osprd_info_t *user_data)
{
    int fd;
    task_lock(task);
    spin_lock(&task->files->file_lock);
    {
#ifdef LINUX_VERSION_CODE <= KERNEL_VERSION(2, 6, 13)
        struct files_struct *f = task->files;
    #else
        struct fdtable *f = task->files->fdt;
    #endif
        for (fd = 0; fd < f->max_fds; fd++)
            if (f->fd[fd])
```

```
                (*callback)(f->fd[fd], user_data);
    }
    spin_unlock(&task->files->file_lock);
    task_unlock(task);
}
```

// Destroy a osprd_info_t.

```
static void cleanup_device(osprd_info_t *d)
{
    wake_up_all(&d->blockq);
    if (d->gd) {
        del_gendisk(d->gd);
        put_disk(d->gd);
    }
    if (d->queue)
        blk_cleanup_queue(d->queue);
    if (d->data)
        vfree(d->data);
}
```

// Initialize a osprd_info_t.

```
static int setup_device(osprd_info_t *d, int which)
{
    memset(d, 0, sizeof(osprd_info_t));

    /* Get memory to store the actual block data. */
    if (!(d->data = vmalloc(nsectors * SECTOR_SIZE)))
        return -1;
    memset(d->data, 0, nsectors * SECTOR_SIZE);

    /* Set up the I/O queue. */
    spin_lock_init(&d->qlock);
    if (!(d->queue = blk_init_queue(osprd_process_request_queue, &d->qlock)))
        return -1;
    blk_queue_hardsect_size(d->queue, SECTOR_SIZE);
    d->queue->queuedata = d;

    /* The gendisk structure. */
    if (!(d->gd = alloc_disk(1)))
        return -1;
    d->gd->major = OSPRD_MAJOR;
    d->gd->first_minor = which;
    d->gd->fops = &osprd_ops;
    d->gd->queue = d->queue;
    d->gd->private_data = d;
    snprintf(d->gd->disk_name, 32, "osprd%c", which + 'a');
    set_capacity(d->gd, nsectors);
    add_disk(d->gd);

    /* Call the setup function. */
    osprd_setup(d);
```

LAB 2

```
    return 0;
}

static void osprd_exit(void);

// The kernel calls this function when the module is loaded.
// It initializes the 4 osprd block devices.

static int __init osprd_init(void)
{
    int i, r;

    // shut up the compiler
    (void) for_each_open_file;
#ifdef osp_spin_lock
    (void) osp_spin_lock;
    (void) osp_spin_unlock;
#endif

    /* Register the block device name. */
    if (register_blkdev(OSPRD_MAJOR, "osprd") < 0) {
        printk(KERN_WARNING "osprd: unable to get major number\n");
        return -EBUSY;
    }

    /* Initialize the device structures. */
    for (i = r = 0; i < NOSPRD; i++)
        if (setup_device(&osprds[i], i) < 0)
            r = -EINVAL;

    if (r < 0) {
        printk(KERN_EMERG "osprd: can't set up device structures\n");
        osprd_exit();
        return -EBUSY;
    } else
        return 0;
}

// The kernel calls this function to unload the osprd module.
// It destroys the osprd devices.

static void osprd_exit(void)
{
    int i;
    for (i = 0; i < NOSPRD; i++)
        cleanup_device(&osprds[i]);
    unregister_blkdev(OSPRD_MAJOR, "osprd");
}

// Tell Linux to call those functions at init and exit time.
module_init(osprd_init);

module_exit(osprd_exit);
```