

LAB 3

```
#include <linux/autofs.h>
#include <linux/version.h>
#ifndef EXPORT_SYMTAB
#define EXPORT_SYMTAB
#endif
#include <linux/module.h>
#include <linux/moduleparam.h>
#include "ospfs.h"
#include <linux/string.h>
#include <linux/slab.h>
#include <linux/file.h>
#include <linux/fs.h>
#include <linux/namei.h>
#include <asm/uaccess.h>
#include <linux/kernel.h>
#include <linux/sched.h>

/*****
 * ospfsmod
 *
 * This is the OSPFS module! It contains both library code for your use,
 * and exercises where you must add code.
 *
 *****/

/* Define eprintk() to be a version of printk(), which prints messages to
 * the console.
 * (If working on a real Linux machine, change KERN_NOTICE to KERN_ALERT or
 * KERN_EMERG so that you are sure to see the messages. By default, the
 * kernel does not print all messages to the console. Levels like KERN_ALERT
 * and KERN_EMERG will make sure that you will see messages.) */
#define eprintk(format, ...) printk(KERN_NOTICE format, ## __VA_ARGS__)

// The actual disk data is just an array of raw memory.
// The initial array is defined in fsmg.c, based on your 'base' directory.
extern uint8_t ospfs_data[];
extern uint32_t ospfs_length;

// A pointer to the superblock; see ospfs.h for details on the struct.
static ospfs_super_t * const ospfs_super =
    (ospfs_super_t *) &ospfs_data[OSPFS_BLKSIZE];

static int change_size(ospfs_inode_t *oi, uint32_t want_size);
static ospfs_dentry_t *find_dentry(ospfs_inode_t *dir_oi, const char *name, int namelen);

/*****
 * FILE SYSTEM OPERATIONS STRUCTURES
 *
 * Linux filesystems are based around three interrelated structures.
 *
 * These are:
 *
 * 1. THE LINUX SUPERBLOCK. This structure represents the whole file system.
 * Example members include the root directory and the number of blocks
```

```
* on the disk.
*
* 2. LINUX INODES. Each file and directory in the file system corresponds
* to an inode. Inode operations include "mkdir" and "create" (add to
* directory).
*
* 3. LINUX FILES. Corresponds to an open file or directory. Operations
* include "read", "write", and "readdir".
*
*
* When Linux wants to perform some file system operation,
* it calls a function pointer provided by the file system type.
* (Thus, Linux file systems are object oriented!)
*
*
* These function pointers are grouped into structures called "operations"
* structures.
*
*
* The initial portion of the file declares all the operations structures we
* need to support ospfsmod: one for the superblock, several for different
* kinds of inodes and files. There are separate inode_operations and
* file_operations structures for OSPFS directories and for regular OSPFS
* files. The structures are actually defined near the bottom of this file.
*/
```

```
// Basic file system type structure
// (links into Linux's list of file systems it supports)
static struct file_system_type ospfs_fs_type;
// Inode and file operations for regular files
static struct inode_operations ospfs_reg_inode_ops;
static struct file_operations ospfs_reg_file_ops;
// Inode and file operations for directories
static struct inode_operations ospfs_dir_inode_ops;
static struct file_operations ospfs_dir_file_ops;
// Inode operations for symbolic links
static struct inode_operations ospfs_symlink_inode_ops;
// Other required operations
static struct dentry_operations ospfs_dentry_ops;
static struct super_operations ospfs_superblock_ops;
```

```
/*****
 * BITVECTOR OPERATIONS
 *
 * OSPFS uses a free bitmap to keep track of free blocks.
 * These bitvector operations, which set, clear, and test individual bits
 * in a bitmap, may be useful.
 */
```

```
// bitvector_set -- Set 'i'th bit of 'vector' to 1.
static inline void
bitvector_set(void *vector, int i)
{
    ((uint32_t *) vector)[i / 32] |= (1 << (i % 32));
}
```

```
// bitvector_clear -- Set 'i'th bit of 'vector' to 0.
static inline void
```

LAB 3

```
bitvector_clear(void *vector, int i)
{
    ((uint32_t *) vector) [i / 32] &= ~(1 << (i % 32));
}

// bitvector_test -- Return the value of the 'i'th bit of 'vector'.
static inline int
bitvector_test(const void *vector, int i)
{
    return (((const uint32_t *) vector) [i / 32] & (1 << (i % 32))) != 0;
}

/*****
 * OSPFS HELPER FUNCTIONS
 */

// ospfs_size2nblocks(size)
// Returns the number of blocks required to hold 'size' bytes of data.
//
// Input:  size -- file size
// Returns: a number of blocks

uint32_t
ospfs_size2nblocks(uint32_t size)
{
    return (size + OSPFS_BLKSIZE - 1) / OSPFS_BLKSIZE;
}

// ospfs_block(blockno)
// Use this function to load a block's contents from "disk".
//
// Input:  blockno -- block number
// Returns: a pointer to that block's data

static void *
ospfs_block(uint32_t blockno)
{
    return &ospfs_data[blockno * OSPFS_BLKSIZE];
}

// ospfs_inode(ino)
// Use this function to load a 'ospfs_inode' structure from "disk".
//
// Input:  ino -- inode number
// Returns: a pointer to the corresponding ospfs_inode structure

static inline ospfs_inode_t *
ospfs_inode(ino_t ino)
{
    ospfs_inode_t *oi;
    if (ino >= ospfs_super->os_ninodes)

        return 0;
    oi = ospfs_block(ospfs_super->os_firstinob);
    return &oi[ino];
}

// ospfs_inode_blockno(oi, offset)
// Use this function to look up the blocks that are part of a file's
// contents.
//
// Inputs:  oi    -- pointer to a OSPFS inode
//          offset -- byte offset into that inode
// Returns: the block number of the block that contains the 'offset'th byte
//          of the file

static inline uint32_t
ospfs_inode_blockno(ospfs_inode_t *oi, uint32_t offset)
{
    uint32_t blockno = offset / OSPFS_BLKSIZE;
    if (offset >= oi->oi_size || oi->oi_ftype == OSPFS_FTYPE_SYMLINK)
        return 0;
    else if (blockno >= OSPFS_NDIRECT + OSPFS_NINDIRECT) {
        uint32_t blockoff = blockno - (OSPFS_NDIRECT + OSPFS_NINDIRECT);
        uint32_t *indirect2_block = ospfs_block(oi->oi_indirect2);
        uint32_t *indirect_block = ospfs_block(indirect2_block[blockoff / OSPFS_NINDIRECT]);
        return indirect_block[blockoff % OSPFS_NINDIRECT];
    } else if (blockno >= OSPFS_NDIRECT) {
        uint32_t *indirect_block = ospfs_block(oi->oi_indirect);
        return indirect_block[blockno - OSPFS_NDIRECT];
    } else
        return oi->oi_direct[blockno];
}

// ospfs_inode_data(oi, offset)
// Use this function to load part of inode's data from "disk",
// where 'offset' is relative to the first byte of inode data.
//
// Inputs:  oi    -- pointer to a OSPFS inode
//          offset -- byte offset into 'oi's data contents
// Returns: a pointer to the 'offset'th byte of 'oi's data contents
//
// Be careful: the returned pointer is only valid within a single block.
// This function is a simple combination of 'ospfs_inode_blockno'
// and 'ospfs_block'.

static inline void *
ospfs_inode_data(ospfs_inode_t *oi, uint32_t offset)
{
    uint32_t blockno = ospfs_inode_blockno(oi, offset);
    return (uint8_t *) ospfs_block(blockno) + (offset % OSPFS_BLKSIZE);
}

/*****
```

LAB 3

* LOW-LEVEL FILE SYSTEM FUNCTIONS

* There are no exercises in this section, and you don't need to understand
* the code.
*/

```
// ospfs_mk_linux_inode(sb, ino)
// Linux's in-memory 'struct inode' structure represents disk
// objects (files and directories). Many file systems have their own
// notion of inodes on disk, and for such file systems, Linux's
// 'struct inode's are like a cache of on-disk inodes.
//
// This function takes an inode number for the OSPFS and constructs
// and returns the corresponding Linux 'struct inode'.
//
// Inputs: sb -- the relevant Linux super_block structure (one per mount)
//         ino -- OSPFS inode number
// Returns: 'struct inode'
```

```
static struct inode *
ospfs_mk_linux_inode(struct super_block *sb, ino_t ino)
```

```
{
    ospfs_inode_t *oi = ospfs_inode(ino);
    struct inode *inode;
```

```
    if (!oi)
        return 0;
    if (!(inode = new_inode(sb)))
        return 0;
```

```
    inode->i_ino = ino;
    // Make it look like everything was created by root.
    inode->i_uid = inode->i_gid = 0;
    inode->i_size = oi->oi_size;
```

```
    if (oi->oi_fstype == OSPFS_FTYPE_REG) {
        // Make an inode for a regular file.
        inode->i_mode = oi->oi_mode | S_IFREG;
        inode->i_op = &ospfs_reg_inode_ops;
        inode->i_fop = &ospfs_reg_file_ops;
        inode->i_nlink = oi->oi_nlink;
```

```
    } else if (oi->oi_fstype == OSPFS_FTYPE_DIR) {
        // Make an inode for a directory.
        inode->i_mode = oi->oi_mode | S_IFDIR;
        inode->i_op = &ospfs_dir_inode_ops;
        inode->i_fop = &ospfs_dir_file_ops;
        inode->i_nlink = oi->oi_nlink + 1 /* dot-dot */;
```

```
    } else if (oi->oi_fstype == OSPFS_FTYPE_SYMLINK) {
        // Make an inode for a symbolic link.
        inode->i_mode = S_IRUSR | S_IRGRP | S_IROTH
            | S_IWUSR | S_IWGRP | S_IWOTH
            | S_IXUSR | S_IXGRP | S_IXOTH | S_IFLNK;
        inode->i_op = &ospfs_symlink_inode_ops;
        inode->i_nlink = oi->oi_nlink;
```

```
    } else
        panic("OSPFS: unknown inode type!");
```

```
    // Access and modification times are now.
    inode->i_mtime = inode->i_atime = inode->i_ctime = CURRENT_TIME;
    return inode;
}
```

```
// ospfs_fill_super, ospfs_get_sb
// These functions are called by Linux when the user mounts a version of
// the OSPFS onto some directory. They help construct a Linux
// 'struct super_block' for that file system.
```

```
static int
ospfs_fill_super(struct super_block *sb, void *data, int flags)
{
    struct inode *root_inode;
```

```
    sb->s_blocksize = OSPFS_BLKSIZE;
    sb->s_blocksize_bits = OSPFS_BLKSIZE_BITS;
    sb->s_magic = OSPFS_MAGIC;
    sb->s_op = &ospfs_superblock_ops;
```

```
    if (!(root_inode = ospfs_mk_linux_inode(sb, OSPFS_ROOT_INO))
        || !(sb->s_root = d_alloc_root(root_inode))) {
        iput(root_inode);
        sb->s_dev = 0;
        return -ENOMEM;
    }
```

```
    return 0;
}
```

```
static int
ospfs_get_sb(struct file_system_type *fs_type, int flags, const char *dev_name, void *data, struct
vfsmount *mount)
{
    return get_sb_single(fs_type, flags, data, ospfs_fill_super, mount);
}
```

```
// ospfs_delete_dentry
// Another bookkeeping function.
```

```
static int
ospfs_delete_dentry(struct dentry *dentry)
{
    return 1;
}
```

```
/* *****
* DIRECTORY OPERATIONS
```

LAB 3

```

*
* EXERCISE: Finish 'ospfs_dir_readdir' and 'ospfs_symlink'.
*/

// ospfs_dir_lookup(dir, dentry, ignore)
// This function implements the "lookup" directory operation, which
// looks up a named entry.
//
// We have written this function for you.
//
// Input: dir -- The Linux 'struct inode' for the directory.
//        You can extract the corresponding 'ospfs_inode_t'
//        by calling 'ospfs_inode' with the relevant inode number.
//        dentry -- The name of the entry being looked up.
// Effect: Looks up the entry named 'dentry'. If found, attaches the
//        entry's 'struct inode' to the 'dentry'. If not found, returns
//        a "negative dentry", which has no inode attachment.

static struct dentry *
ospfs_dir_lookup(struct inode *dir, struct dentry *dentry, struct nameidata *ignore)
{
    // Find the OSPFS inode corresponding to 'dir'
    ospfs_inode_t *dir_oi = ospfs_inode(dir->i_ino);
    struct inode *entry_inode = NULL;
    int entry_off;

    // Make sure filename is not too long
    if (dentry->d_name.len > OSPFS_MAXNAMELEN)
        return (struct dentry *) ERR_PTR(-ENAMETOOLONG);

    // Mark with our operations
    dentry->d_op = &ospfs_dentry_ops;

    // Search through the directory block
    for (entry_off = 0; entry_off < dir_oi->oi_size;
         entry_off += OSPFS_DIRENTRY_SIZE) {
        // Find the OSPFS inode for the entry
        ospfs_direntry_t *od = ospfs_inode_data(dir_oi, entry_off);

        // Set 'entry_inode' if we find the file we are looking for
        if (od->od_ino > 0
            && strlen(od->od_name) == dentry->d_name.len
            && memcmp(od->od_name, dentry->d_name.name, dentry->d_name.len) == 0) {
            entry_inode = ospfs_mk_linux_inode(dir->i_sb, od->od_ino);
            if (!entry_inode)
                return (struct dentry *) ERR_PTR(-EINVAL);
            break;
        }
    }

    // We return a dentry whether or not the file existed.
    // The file exists if and only if 'entry_inode != NULL'.
    // If the file doesn't exist, the dentry is called a "negative dentry".

    // d_splice_alias() attaches the inode to the dentry.

    // If it returns a new dentry, we need to set its operations.
    if ((dentry = d_splice_alias(entry_inode, dentry)))
        dentry->d_op = &ospfs_dentry_ops;
    return dentry;
}

// ospfs_dir_readdir(filp, dirent, filldir)
// This function is called when the kernel reads the contents of a directory
// (i.e. when file_operations.readdir is called for the inode).
//
// Inputs: filp -- The 'struct file' structure corresponding to
//        the open directory.
//        The most important member is 'filp->f_pos', the
//        File POSition. This remembers how far into the
//        directory we are, so if the user calls 'readdir'
//        twice, we don't forget our position.
//        This function must update 'filp->f_pos'.
//        dirent -- Used to pass to 'filldir'.
//        filldir -- A pointer to a callback function.
//        This function should call 'filldir' once for each
//        directory entry, passing it six arguments:
//        (1) 'dirent'.
//        (2) The directory entry's name.
//        (3) The length of the directory entry's name.
//        (4) The 'f_pos' value corresponding to the directory entry.
//        (5) The directory entry's inode number.
//        (6) DT_REG, for regular files; DT_DIR, for subdirectories;
//        or DT_LNK, for symbolic links.
//        This function should stop returning directory
//        entries either when the directory is complete, or
//        when 'filldir' returns < 0, whichever comes first.
//
// Returns: 1 at end of directory, 0 if filldir returns < 0 before the end
//        of the directory, and -(error number) on error.
//
// EXERCISE: Finish implementing this function.

static int
ospfs_dir_readdir(struct file *filp, void *dirent, filldir_t filldir)
{
    struct inode *dir_inode = filp->f_dentry->d_inode;
    ospfs_inode_t *dir_oi = ospfs_inode(dir_inode->i_ino);
    uint32_t f_pos = filp->f_pos;
    int r = 0; /* Error return value, if any */
    int ok_so_far = 0; /* Return value from 'filldir' */

    // f_pos is an offset into the directory's data, plus two.
    // The "plus two" is to account for "." and "..".
    if (r == 0 && f_pos == 0) {
        ok_so_far = filldir(dirent, ".", 1, f_pos, dir_inode->i_ino, DT_DIR);
        if (ok_so_far >= 0)
            f_pos++;
    }
}

```

LAB 3

```

if (r == 0 && ok_so_far >= 0 && f_pos == 1) {
    ok_so_far = filldir(dirent, "..", 2, f_pos, filp->f_dentry->d_parent->d_inode->i_ino, DT_DIR);
    if (ok_so_far >= 0)
        f_pos++;
}

while (r == 0 && ok_so_far >= 0 && f_pos >= 2) {
    ospfs_dirent_t *od;
    ospfs_inode_t *entry_oi;

    if (f_pos > dir_oi->oi_size * OSPFS_DIRENTRY_SIZE) { /* TODO: error cond */
        r = 1;
        break;
    }

    /* Get a pointer to the next entry (od) in the directory.
     * The file system interprets the contents of a
     * directory-file as a sequence of ospfs_dirent structures.
     * You will find 'f_pos' and 'ospfs_inode_data' useful.
     *
     * Then use the fields of that file to fill in the directory
     * entry. To figure out whether a file is a regular file or
     * another directory, use 'ospfs_inode' to get the directory
     * entry's corresponding inode, and check out its 'oi_fstype'
     * member.
     *
     * Make sure you ignore blank directory entries! (Which have
     * an inode number of 0.)
     *
     * If the current entry is successfully read (the call to
     * filldir returns >= 0), or the current entry is skipped,
     * your function should advance f_pos by the proper amount to
     * advance to the next directory entry.
     */

    od = ospfs_inode_data(dir_oi, f_pos * OSPFS_DIRENTRY_SIZE);
    entry_oi = ospfs_inode(od->od_ino);

    if(entry_oi != 0) {
        switch(entry_oi->oi_fstype) {
            case OSPFS_FSTYPE_REG:
                ok_so_far = filldir(dirent, od->od_name, strlen(od->od_name), f_pos, od->od_ino, DT_REG);
                break;
            case OSPFS_FSTYPE_DIR:
                ok_so_far = filldir(dirent, od->od_name, strlen(od->od_name), f_pos, od->od_ino, DT_DIR);
                break;
            case OSPFS_FSTYPE_SYMLINK:
                ok_so_far = filldir(dirent, od->od_name, strlen(od->od_name), f_pos, od->od_ino, DT_LNK);
                break;
            default: {
                r=1;
                continue;
            }
        }
    }
}

```

```

    }

    f_pos++;
}

filp->f_pos = f_pos;
return r;
}

// ospfs_unlink(dirino, dentry)
// This function is called to remove a file.
//
// Inputs: dirino -- You may ignore this.
//          dentry -- The 'struct dentry' structure, which contains the inode
//                   the directory entry points to and the directory entry's
//                   directory.
//
// Returns: 0 if success and -ENOENT on entry not found.
//
// EXERCISE: Make sure that deleting symbolic links works correctly.

static int
ospfs_unlink(struct inode *dirino, struct dentry *dentry)
{
    ospfs_inode_t *oi = ospfs_inode(dentry->d_inode->i_ino);
    ospfs_inode_t *dir_oi = ospfs_inode(dentry->d_parent->d_inode->i_ino);
    int entry_off;
    ospfs_dirent_t *od;

    od = NULL;
    for (entry_off = 0; entry_off < dir_oi->oi_size;
         entry_off += OSPFS_DIRENTRY_SIZE) {
        od = ospfs_inode_data(dir_oi, entry_off);
        if (od->od_ino > 0
            && strlen(od->od_name) == dentry->d_name.len
            && memcmp(od->od_name, dentry->d_name.name, dentry->d_name.len) == 0)
            break;
    }

    if (entry_off == dir_oi->oi_size) {
        return -ENOENT;
    }

    od->od_ino = 0;
    oi->oi_nlink--;

    if (oi->oi_fstype != OSPFS_FSTYPE_SYMLINK && oi->oi_nlink == 0)
        return change_size(oi, 0);

    return 0;
}

```

LAB 3

```
/******  
 * FREE-BLOCK BITMAP OPERATIONS  
 *  
 * EXERCISE: Implement these functions.  
 */  
  
// allocate_block()  
// Use this function to allocate a block.  
//  
// Inputs: none  
// Returns: block number of the allocated block,  
//         or 0 if the disk is full  
//  
// This function searches the free-block bitmap, which starts at Block 2, for  
// a free block, allocates it (by marking it non-free), and returns the block  
// number to the caller. The block itself is not touched.  
//  
// Note: A value of 0 for a bit indicates the corresponding block is  
//        allocated; a value of 1 indicates the corresponding block is free.  
//  
// You can use the functions bitvector_set(), bitvector_clear(), and  
// bitvector_test() to do bit operations on the map.  
  
static uint32_t  
allocate_block(void)  
{  
    void *bitmap = ospfs_block(OSPFS_FREEMAP_BLK);  
  
    int i;  
    for (i = 0; i < ospfs_super->os_nblocks; i++) {  
        if (bitvector_test(bitmap, i)) {  
            bitvector_clear(bitmap, i);  
            return i;  
        }  
    }  
  
    return 0;  
}  
  
// free_block(blockno)  
// Use this function to free an allocated block.  
//  
// Inputs: blockno -- the block number to be freed  
// Returns: none  
//  
// This function should mark the named block as free in the free-block  
// bitmap. (You might want to program defensively and make sure the block  
// number isn't obviously bogus: the boot sector, superblock, free-block  
// bitmap, and inode blocks must never be freed. But this is not required.)  
  
static void  
free_block(uint32_t blockno)  
{  
    void *bitmap = ospfs_block(OSPFS_FREEMAP_BLK);
```

```
    bitvector_set(bitmap, blockno);  
}
```

```
/******  
 * FILE OPERATIONS  
 *  
 * EXERCISE: Finish off change_size, read, and write.  
 *  
 * The find_*, add_block, and remove_block functions are only there to support  
 * the change_size function. If you prefer to code change_size a different  
 * way, then you may not need these functions.  
 *  
 */
```

```
// The following functions are used in our code to unpack a block number into  
// its constituent pieces: the doubly indirect block number (if any), the  
// indirect block number (which might be one of many in the doubly indirect  
// block), and the direct block number (which might be one of many in an  
// indirect block). We use these functions in our implementation of  
// change_size.
```

```
// int32_t indir2_index(uint32_t b)  
// Returns the doubly-indirect block index for file block b.  
//  
// Inputs: b -- the zero-based index of the file block (e.g., 0 for the first  
//          block, 1 for the second, etc.)  
// Returns: 0 if block index 'b' requires using the doubly indirect  
//          block, -1 if it does not.  
//  
// EXERCISE: Fill in this function.
```

```
static int32_t  
indir2_index(uint32_t b)  
{  
    if (b < OSPFS_NDIRECT + OSPFS_NINDIRECT) {  
        return -1;  
    }  
    else {  
        return 0;  
    }  
}
```

```
// int32_t indir_index(uint32_t b)  
// Returns the indirect block index for file block b.  
//  
// Inputs: b -- the zero-based index of the file block  
// Returns: -1 if b is one of the file's direct blocks;  
//          0 if b is located under the file's first indirect block;  
//          otherwise, the offset of the relevant indirect block within  
//          the doubly indirect block.  
//  
// EXERCISE: Fill in this function.
```

LAB 3

```
static int32_t
indir_index(uint32_t b)
{
    if (b < OSPFS_NDIRECT) {
        return -1;
    }
    else if (b < OSPFS_NDIRECT + OSPFS_NINDIRECT) {
        return 0;
    }
    else {
        return (b - OSPFS_NDIRECT - OSPFS_NINDIRECT) / OSPFS_NINDIRECT;
    }
}
```

```
// int32_t indir_index(uint32_t b)
// Returns the indirect block index for file block b.
//
// Inputs: b -- the zero-based index of the file block
// Returns: the index of block b in the relevant indirect block or the direct
// block array.
//
// EXERCISE: Fill in this function.
```

```
static int32_t
direct_index(uint32_t b)
{
    if (b < OSPFS_NDIRECT) {
        return b;
    }
    else if (b < OSPFS_NDIRECT + OSPFS_NINDIRECT) {
        return b - OSPFS_NDIRECT;
    }
    else {
        return (b - OSPFS_NDIRECT) % OSPFS_NINDIRECT;
    }
}
```

```
// add_block(ospfs_inode_t *oi)
// Adds a single data block to a file, adding indirect and
// doubly-indirect blocks if necessary. (Helper function for
// change_size).
//
// Inputs: oi -- pointer to the file we want to grow
// Returns: 0 if successful, < 0 on error. Specifically:
// -ENOSPC if you are unable to allocate a block
// -EIO for any other error.
// If the function is successful, then oi->oi_size
// should be set to the maximum file size in bytes that could
// fit in oi's data blocks. If the function returns an error,
// then oi->oi_size should remain unchanged. Any newly
// allocated blocks should be erased (set to zero).
```

```
//
// EXERCISE: Finish off this function.
//
// Remember that allocating a new data block may require allocating
// as many as three disk blocks, depending on whether a new indirect
// block and/or a new indirect^2 block is required. If the function
// fails with -ENOSPC or -EIO, then you need to make sure that you
// free any indirect (or indirect^2) blocks you may have allocated!
//
// Also, make sure you:
// 1) zero out any new blocks that you allocate
// 2) store the disk block number of any newly allocated block
// in the appropriate place in the inode or one of the
// indirect blocks.
// 3) update the oi->oi_size field
```

```
static int
add_block(ospfs_inode_t *oi)
{
    // current number of blocks in file
    uint32_t n = ospfs_size2nblocks(oi->oi_size);

    // keep track of allocations to free in case of -ENOSPC
    uint32_t allocated[2] = { 0, 0 };
    uint32_t direct, indir;

    if (n < 0)
        return -EIO;
    else if (n < OSPFS_NDIRECT) {
        direct = allocate_block();
        if (!direct)
            return -ENOSPC;
        memset(ospfs_block(direct), 0, OSPFS_BLKSIZE);

        oi->oi_direct[n] = direct;
    }
    else if (n < OSPFS_NDIRECT + OSPFS_NINDIRECT) {
        if (!oi->oi_indirect) {
            allocated[0] = allocate_block();
            if (!allocated[0])
                return -ENOSPC;
            memset(ospfs_block(allocated[0]), 0, OSPFS_BLKSIZE);

            oi->oi_indirect = allocated[0];
        }

        direct = allocate_block();
        if (!direct) {
            if (allocated[0]) {
                free_block(allocated[0]);
                oi->oi_indirect = 0;
            }
            return -ENOSPC;
        }
        memset(ospfs_block(direct), 0, OSPFS_BLKSIZE);
```

LAB 3

```

    ((uint32_t*) ospfs_block(oi->oi_indirect))[direct_index(n)] = direct;
}
else if (n < OSPFS_MAXFILEBLKS) {
    if (!oi->oi_indirect2) {
        allocated[0] = allocate_block();
        if (!allocated[0])
            return -ENOSPC;
        memset(ospfs_block(allocated[0]), 0, OSPFS_BLKSIZE);

        oi->oi_indirect2 = allocated[0];
    }

    indir = ((uint32_t *) ospfs_block(oi->oi_indirect2))[indir_index(n)];

    if (!indir) {
        allocated[1] = allocate_block();
        if (!allocated[1]) {
            if (allocated[0])
                free_block(allocated[0]);
            return -ENOSPC;
        }
        memset(ospfs_block(allocated[1]), 0, OSPFS_BLKSIZE);

        indir = allocated[1];
    }

    direct = allocate_block();
    if (!direct) {
        if (allocated[0]) {
            free_block(allocated[0]);
            oi->oi_indirect2 = 0;
        }
        if (allocated[1])
            free_block(allocated[1]);
        return -ENOSPC;
    }
    memset(ospfs_block(direct), 0, OSPFS_BLKSIZE);

    ((uint32_t *) ospfs_block(indir))[direct_index(n)] = direct;
}
else {
    return -ENOSPC;
}

oi->oi_size += OSPFS_BLKSIZE;

return 0;
}

// remove_block(ospfs_inode_t *oi)
// Removes a single data block from the end of a file, freeing
// any indirect and indirect^2 blocks that are no

```

```

// longer needed. (Helper function for change_size)
//
// Inputs: oi -- pointer to the file we want to shrink
// Returns: 0 if successful, < 0 on error.
//     If the function is successful, then oi->oi_size
//     should be set to the maximum file size that could
//     fit in oi's blocks. If the function returns -EIO (for
//     instance if an indirect block that should be there isn't),
//     then oi->oi_size should remain unchanged.
//
// EXERCISE: Finish off this function.
//
// Remember that you must free any indirect and doubly-indirect blocks
// that are no longer necessary after shrinking the file. Removing a
// single data block could result in as many as 3 disk blocks being
// deallocated. Also, if you free a block, make sure that
// you set the block pointer to 0. Don't leave pointers to
// deallocated blocks laying around!

static int
remove_block(ospfs_inode_t *oi)
{
    // current number of blocks in file
    uint32_t n = ospfs_size2nblocks(oi->oi_size);

    if (n < 0) {
        return -EIO;
    }
    else if (n == 0) {
        return 0;
    }

    n = n - 1;

    if (n < OSPFS_NDIRECT) {
        free_block(oi->oi_direct[n]);
        oi->oi_direct[n] = 0;
    }
    else if (n < OSPFS_NDIRECT + OSPFS_NINDIRECT) {
        uint32_t* indir = (uint32_t *) ospfs_block(oi->oi_indirect);

        free_block(indir[direct_index(n)]);
        indir[direct_index(n)] = 0;

        if (direct_index(n) == 0) {
            free_block(oi->oi_indirect);
            oi->oi_indirect = 0;
        }
    }
    else if (n < OSPFS_MAXFILEBLKS) {
        uint32_t* indir2 = (uint32_t *) ospfs_block(oi->oi_indirect2);
        uint32_t* indir = (uint32_t *) ospfs_block(indir2[indir_index(n)]);

        free_block(indir[direct_index(n)]);
        indir[direct_index(n)] = 0;
    }
}

```


LAB 3

```
if (direct_index(n) == 0) {
    free_block(indir2[indir_index(n)]);
    indir2[indir_index(n)] = 0;
}

if (indir_index(n) == 0) {
    free_block(oi->oi_indirect2);
    oi->oi_indirect2 = 0;
}
}
else
    return -EIO;

oi->oi_size -= OSPFS_BLKSIZE;

return 0;
}

// change_size(oi, want_size)
// Use this function to change a file's size, allocating and freeing
// blocks as necessary.
//
// Inputs: oi-- pointer to the file whose size we're changing
//         want_size -- the requested size in bytes
// Returns: 0 on success, < 0 on error. In particular:
// -ENOSPC: if there are no free blocks available
// -EIO:    an I/O error -- for example an indirect block should
//         exist, but doesn't
// If the function succeeds, the file's oi_size member should be
// changed to want_size, with blocks allocated as appropriate.
// Any newly-allocated blocks should be erased (set to 0).
// If there is an -ENOSPC error when growing a file,
// the file size and allocated blocks should not change from their
// original values!!!
// (However, if there is an -EIO error, do not worry too much about
// restoring the file.)
//
// If want_size has the same number of blocks as the current file, life
// is good -- the function is pretty easy. But the function might have
// to add or remove blocks.
//
// If you need to grow the file, then do so by adding one block at a time
// using the add_block function you coded above. If one of these additions
// fails with -ENOSPC, you must shrink the file back to its original size!
//
// If you need to shrink the file, remove blocks from the end of
// the file one at a time using the remove_block function you coded above.
//
// Also: Don't forget to change the size field in the metadata of the file.
// (The value that the final add_block or remove_block set it to
// is probably not correct).
//
// EXERCISE: Finish off this function.
```

```
static int
change_size(ospfs_inode_t *oi, uint32_t new_size)
{
    uint32_t old_size = oi->oi_size;
    int r = 0;

    while (ospfs_size2nblocks(oi->oi_size) < ospfs_size2nblocks(new_size)) {
        r = add_block(oi);
        if (r < 0)
            break;
    }

    if (r == -EIO)
        return -EIO;
    else if (r == -ENOSPC) {
        while (ospfs_size2nblocks(oi->oi_size) > ospfs_size2nblocks(old_size))
            r = remove_block(oi);

        oi->oi_size = old_size;
        return -ENOSPC;
    }

    while (ospfs_size2nblocks(oi->oi_size) > ospfs_size2nblocks(new_size)) {
        r = remove_block(oi);
        if (r < 0)
            return r;
    }

    oi->oi_size = new_size;

    return 0;
}

// ospfs_notify_change
// This function gets called when the user changes a file's size,
// owner, or permissions, among other things.
// OSPFS only pays attention to file size changes (see change_size above).
// We have written this function for you -- except for file quotas.

static int
ospfs_notify_change(struct dentry *dentry, struct iattr *attr)
{
    struct inode *inode = dentry->d_inode;
    ospfs_inode_t *oi = ospfs_inode(inode->i_ino);
    int retval = 0;

    if (attr->ia_valid & ATTR_SIZE) {
        if (oi->oi_ftype == OSPFS_FTYPE_DIR)
            return -EPERM;
        if ((retval = change_size(oi, attr->ia_size)) < 0)
            goto out;
    }
}
```

LAB 3

```
if (attr->ia_valid & ATTR_MODE)
    oi->oi_mode = attr->ia_mode;

if ((retval = inode_change_ok(inode, attr)) < 0
    || (retval = inode_setattr(inode, attr)) < 0)
    goto out;

out:
return retval;
}

// ospfs_read
// Linux calls this function to read data from a file.
// It is the file_operations.read callback.
//
// Inputs: filp -- a file pointer
//         buffer -- a user space ptr where data should be copied
//         count -- the amount of data requested
//         f_pos -- points to the file position
// Returns: Number of chars read on success, -(error code) on error.
//
// This function copies the corresponding bytes from the file into the user
// space ptr (buffer). Use copy_to_user() to accomplish this.
// The current file position is passed into the function
// as 'f_pos'; read data starting at that position, and update the position
// when you're done.
//
// EXERCISE: Complete this function.

static ssize_t
ospfs_read(struct file *filp, char __user *buffer, size_t count, loff_t *f_pos)
{
    ospfs_inode_t *oi = ospfs_inode(filp->f_dentry->d_inode->i_ino);
    int retval = 0;
    size_t amount = 0;

    // Make sure we don't read past the end of the file!
    // Change 'count' so we never read past the end of the file.
    if (*f_pos + count > oi->oi_size)
        count = oi->oi_size - *f_pos;

    // Copy the data to user block by block
    while (amount < count && retval >= 0) {
        uint32_t blockno = ospfs_inode_blockno(oi, *f_pos);
        uint32_t n;
        char *data;

        if (blockno == 0) {
            retval = -EIO;
            goto done;
        }

        data = ospfs_block(blockno);

        // Figure out how much data is left in this block to read.
        // Copy data into user space. Return -EFAULT if unable to write
        // into user space.
        // Use variable 'n' to track number of bytes moved.
        n = (count + (*f_pos % OSPFS_BLKSIZE) - amount > OSPFS_BLKSIZE ?
            OSPFS_BLKSIZE - (*f_pos % OSPFS_BLKSIZE) : count - amount);
        retval = copy_to_user(buffer, data, n);

        if (retval < 0) {
            retval = -EFAULT;
            goto done;
        }

        buffer += n;
        amount += n;
        *f_pos += n;
    }

    done:
    return (retval >= 0 ? amount : retval);
}

// ospfs_write
// Linux calls this function to write data to a file.
// It is the file_operations.write callback.
//
// Inputs: filp -- a file pointer
//         buffer -- a user space ptr where data should be copied from
//         count -- the amount of data to write
//         f_pos -- points to the file position
// Returns: Number of chars written on success, -(error code) on error.
//
// This function copies the corresponding bytes from the user space ptr
// into the file. Use copy_from_user() to accomplish this. Unlike read(),
// where you cannot read past the end of the file, it is OK to write past
// the end of the file; this should simply change the file's size.
//
// EXERCISE: Complete this function.

static ssize_t
ospfs_write(struct file *filp, const char __user *buffer, size_t count, loff_t *f_pos)
{
    ospfs_inode_t *oi = ospfs_inode(filp->f_dentry->d_inode->i_ino);
    int retval = 0;
    size_t amount = 0;

    // Support files opened with the O_APPEND flag. To detect O_APPEND,
    // use struct file's f_flags field and the O_APPEND bit.
    if (filp->f_flags & O_APPEND)
        *f_pos = oi->oi_size;

    if ((*f_pos + count) > oi->oi_size)
        if (change_size(oi, (*f_pos + count)) < 0)
            goto done;
}
```

LAB 3

```
// Copy data block by block
while (amount < count && retval >= 0) {
    uint32_t blockno = ospfs_inode_blockno(oi, *f_pos);
    uint32_t n;
    char *data;

    if (blockno == 0) {
        retval = -EIO;
        goto done;
    }

    data = ospfs_block(blockno);

    // Figure out how much data is left in this block to write.
    // Copy data from user space. Return -EFAULT if unable to read
    // read user space.
    // Keep track of the number of bytes moved in 'n'.
    n = OSPFS_BLKSIZE - (*f_pos % OSPFS_BLKSIZE);

    if (n > count - amount)
        n = count - amount;

    if (copy_from_user(data + (*f_pos % OSPFS_BLKSIZE), buffer, n) != 0)
        return -EFAULT;

    buffer += n;
    amount += n;
    *f_pos += n;
}

done:
return (retval >= 0 ? amount : retval);
}

// find_dirent(dir_oi, name, namelen)
// Looks through the directory to find an entry with name 'name' (length
// in characters 'namelen'). Returns a pointer to the directory entry,
// if one exists, or NULL if one does not.
//
// Inputs: dir_oi -- the OSP inode for the directory
// name -- name to search for
// namelen -- length of 'name'. (If -1, then use strlen(name).)
//
// We have written this function for you.

static ospfs_dirent_t *
find_dirent(ospfs_inode_t *dir_oi, const char *name, int namelen)
{
    int off;
    if (namelen < 0)
        namelen = strlen(name);
    for (off = 0; off < dir_oi->oi_size; off += OSPFS_DIRENTRY_SIZE) {
        ospfs_dirent_t *od = ospfs_inode_data(dir_oi, off);
```

```
        if (od->od_ino
            && strlen(od->od_name) == namelen
            && memcmp(od->od_name, name, namelen) == 0)
            return od;
    }
    return 0;
}
```

```
// create_blank_dirent(dir_oi)
// 'dir_oi' is an OSP inode for a directory.
// Return a blank directory entry in that directory. This might require
// adding a new block to the directory. Returns an error pointer (see
// below) on failure.
//
// ERROR POINTERS: The Linux kernel uses a special convention for returning
// error values in the form of pointers. Here's how it works.
// - ERR_PTR(errno): Creates a pointer value corresponding to an error.
// - IS_ERR(ptr): Returns true iff 'ptr' is an error value.
// - PTR_ERR(ptr): Returns the error value for an error pointer.
// For example:
//
// static ospfs_dirent_t *create_blank_dirent(...) {
//     return ERR_PTR(-ENOSPC);
// }
// static int ospfs_create(...) {
//     ...
//     ospfs_dirent_t *od = create_blank_dirent(...);
//     if (IS_ERR(od))
//         return PTR_ERR(od);
//     ...
// }
//
// The create_blank_dirent function should use this convention.
//
// EXERCISE: Write this function.
```

```
static ospfs_dirent_t *
create_blank_dirent(ospfs_inode_t *dir_oi)
{
    ospfs_dirent_t *dir_entry;
    int retval;

    uint32_t offset;
    for (offset = 0; offset < dir_oi->oi_size; offset += OSPFS_DIRENTRY_SIZE) {
        dir_entry = ospfs_inode_data(dir_oi, offset);
        if (dir_entry->od_ino == 0)
            return dir_entry;
    }

    retval = add_block(dir_oi);
    if (retval < 0)
        return ERR_PTR(retval);

    dir_entry = ospfs_inode_data(dir_oi, offset);
```

LAB 3

```

    return dir_entry;
}

// ospfs_link(src_dentry, dir, dst_dentry)
// Linux calls this function to create hard links.
// It is the ospfs_dir_inode_ops.link callback.
//
// Inputs: src_dentry -- a pointer to the dentry for the source file. This
//         file's inode contains the real data for the hard
//         linked filae. The important elements are:
//         src_dentry->d_name.name
//         src_dentry->d_name.len
//         src_dentry->d_inode->i_ino
// dir     -- a pointer to the containing directory for the new
//         hard link.
// dst_dentry -- a pointer to the dentry for the new hard link file.
//         The important elements are:
//         dst_dentry->d_name.name
//         dst_dentry->d_name.len
//         dst_dentry->d_inode->i_ino
//         Two of these values are already set. One must be
//         set by you, which one?
// Returns: 0 on success, -(error code) on error. In particular:
//         -ENAMETOOLONG if dst_dentry->d_name.len is too large, or
//         'symname' is too long;
//         -EEXIST      if a file named the same as 'dst_dentry' already
//                     exists in the given 'dir';
//         -ENOSPC      if the disk is full & the file can't be created;
//         -EIO         on I/O error.
//
// EXERCISE: Complete this function.

static int
ospfs_link(struct dentry *src_dentry, struct inode *dir, struct dentry *dst_dentry) {
    ospfs_dirent_t *link;

    if (dst_dentry->d_name.len > OSPFS_MAXNAMELEN)
        return -ENAMETOOLONG;
    if (find_dirent(ospfs_inode(dir->i_ino),
        dst_dentry->d_name.name, dst_dentry->d_name.len))
        return -EEXIST;

    link = create_blank_dirent(ospfs_inode(dir->i_ino));
    if (IS_ERR(link))
        return PTR_ERR(link);

    link->od_ino = src_dentry->d_inode->i_ino;
    memcpy(link->od_name, dst_dentry->d_name.name, dst_dentry->d_name.len);
    link->od_name[dst_dentry->d_name.len] = '\0';

    ospfs_inode(src_dentry->d_inode->i_ino)->oi_nlink++;

    return 0;
}

```

```

// ospfs_create
// Linux calls this function to create a regular file.
// It is the ospfs_dir_inode_ops.create callback.
//
// Inputs: dir -- a pointer to the containing directory's inode
//         dentry -- the name of the file that should be created
//         The only important elements are:
//         dentry->d_name.name: filename (char array, not null
//         terminated)
//         dentry->d_name.len: length of filename
//         mode -- the permissions mode for the file (set the new
//         inode's oi_mode field to this value)
//         nd -- ignore this
// Returns: 0 on success, -(error code) on error. In particular:
//         -ENAMETOOLONG if dentry->d_name.len is too large;
//         -EEXIST      if a file named the same as 'dentry' already
//                     exists in the given 'dir';
//         -ENOSPC      if the disk is full & the file can't be created;
//         -EIO         on I/O error.
//
// We have provided strictly less skeleton code for this function than for
// the others. Here's a brief outline of what you need to do:
// 1. Check for the -EEXIST error and find an empty directory entry using the
//    helper functions above.
// 2. Find an empty inode. Set the 'entry_ino' variable to its inode number.
// 3. Initialize the directory entry and inode.
//
// EXERCISE: Complete this function.

static int
ospfs_create(struct inode *dir, struct dentry *dentry, int mode, struct nameidata *nd)
{
    ospfs_inode_t *dir_oi = ospfs_inode(dir->i_ino);
    ospfs_dirent_t *dir_entry;
    ospfs_inode_t *inode;
    uint32_t entry_ino = 0;

    if (dentry->d_name.len > OSPFS_MAXNAMELEN)
        return -ENAMETOOLONG;
    if (find_dirent(dir_oi, dentry->d_name.name, dentry->d_name.len))
        return -EEXIST;

    dir_entry = create_blank_dirent(dir_oi);
    if (IS_ERR(dir_entry))
        return PTR_ERR(dir_entry);

    for (entry_ino = 0; entry_ino < ospfs_super->os_ninodes; entry_ino++) {
        inode = ospfs_inode(entry_ino);
        if (inode->oi_nlink == 0)
            break;
    }
    if (entry_ino == ospfs_super->os_ninodes)
        return -ENOSPC;

    dir_entry->od_ino = entry_ino;

```

LAB 3

```

memcpy(dir_entry->od_name, dentry->d_name.name, dentry->d_name.len);
dir_entry->od_name[dentry->d_name.len] = '\0';
// initialize file
inode->oi_size = 0;
inode->oi_ftype = OSPFS_FTYPE_REG;
inode->oi_mode = mode;
inode->oi_nlink = 1;

/* Execute this code after your function has successfully created the
file. Set entry_ino to the created file's inode number before
getting here. */
{
    struct inode *i = ospfs_mk_linux_inode(dir->i_sb, entry_ino);
    if (!i)
        return -ENOMEM;
    d_instantiate(dentry, i);
    return 0;
}

// ospfs_symlink(dirino, dentry, symname)
// Linux calls this function to create a symbolic link.
// It is the ospfs_dir_inode_ops.symlink callback.
//
// Inputs: dir -- a pointer to the containing directory's inode
//         dentry -- the name of the file that should be created
//         The only important elements are:
//         dentry->d_name.name: filename (char array, not null
//         terminated)
//         dentry->d_name.len: length of filename
//         symname -- the symbolic link's destination
//
// Returns: 0 on success, -(error code) on error. In particular:
//         -ENAMETOOLONG if dentry->d_name.len is too large, or
//         'symname' is too long;
//         -EEXIST if a file named the same as 'dentry' already
//         exists in the given 'dir';
//         -ENOSPC if the disk is full & the file can't be created;
//         -EIO on I/O error.
//
// EXERCISE: Complete this function.

static int
ospfs_symlink(struct inode *dir, struct dentry *dentry, const char *symname)
{
    ospfs_inode_t *dir_oi = ospfs_inode(dir->i_ino);
    uint32_t entry_ino = 0;
    ospfs_symlink_inode_t *link;

    if (dentry->d_name.len > OSPFS_MAXNAMELEN ||
        strlen(symname) > OSPFS_MAXSYMLINKLEN)
        return -ENAMETOOLONG;
    if (find_dirent(ospfs_inode(dir->i_ino),
        dentry->d_name.name, dentry->d_name.len))

```

```

        return -EEXIST;

    entry_ino = ospfs_create(dir, dentry, dir_oi->oi_mode, NULL);
    if (entry_ino < 0)
        return entry_ino;
    entry_ino = find_dirent(ospfs_inode(dir->i_ino),
        dentry->d_name.name, dentry->d_name.len)->od_ino;
    link = (ospfs_symlink_inode_t *) ospfs_inode(entry_ino);

    link->oi_size = strlen(symname);
    link->oi_ftype = OSPFS_FTYPE_SYMLINK;
    link->oi_nlink = 1;
    memcpy(link->oi_symlink, symname, strlen(symname));

    /* Execute this code after your function has successfully created the
file. Set entry_ino to the created file's inode number before
getting here. */
    {
        struct inode *i = ospfs_mk_linux_inode(dir->i_sb, entry_ino);
        if (!i)
            return -ENOMEM;
        d_instantiate(dentry, i);
        return 0;
    }
}

// ospfs_follow_link(dentry, nd)
// Linux calls this function to follow a symbolic link.
// It is the ospfs_symlink_inode_ops.follow_link callback.
//
// Inputs: dentry -- the symbolic link's directory entry
//         nd -- to be filled in with the symbolic link's destination
//
// Exercise: Expand this function to handle conditional symlinks. Conditional
// symlinks will always be created by users in the following form
// root?/path/1:/path/2.
// (hint: Should the given form be changed in any way to make this method
// easier? With which character do most functions expect C strings to end?)
//

static void *
ospfs_follow_link(struct dentry *dentry, struct nameidata *nd)
{
    ospfs_symlink_inode_t *oi =
        (ospfs_symlink_inode_t *) ospfs_inode(dentry->d_inode->i_ino);

    if (strcmp(oi->oi_symlink, "root?") == 0) {
        int pivot = strchr(oi->oi_symlink, ':') - oi->oi_symlink;

        if (current->uid == 0) {
            oi->oi_symlink[pivot] = '\0';
            nd_set_link(nd, oi->oi_symlink + 5 + 1);
        }
        else

```

LAB 3

```
        nd_set_link(nd, oi->oi_symlink + pivot + 1);
    }
    else
        nd_set_link(nd, oi->oi_symlink);

    return (void *) 0;
}
```

// Define the file system operations structures mentioned above.

```
static struct file_system_type ospfs_fs_type = {
    .owner      = THIS_MODULE,
    .name       = "ospfs",
    .get_sb     = ospfs_get_sb,
    .kill_sb    = kill_anon_super
};
```

```
static struct inode_operations ospfs_reg_inode_ops = {
    .setattr= ospfs_notify_change
};
```

```
static struct file_operations ospfs_reg_file_ops = {
    .llseek   = generic_file_llseek,
    .read     = ospfs_read,
    .write    = ospfs_write
};
```

```
static struct inode_operations ospfs_dir_inode_ops = {
    .lookup   = ospfs_dir_lookup,
    .link     = ospfs_link,
    .unlink   = ospfs_unlink,
    .create   = ospfs_create,
    .symlink  = ospfs_symlink
};
```

```
static struct file_operations ospfs_dir_file_ops = {
    .read     = generic_read_dir,
    .readdir  = ospfs_dir_readdir
};
```

```
static struct inode_operations ospfs_symlink_inode_ops = {
    .readlink= generic_readlink,
    .follow_link = ospfs_follow_link
};
```

```
static struct dentry_operations ospfs_dentry_ops = {
    .d_delete = ospfs_delete_dentry
};
```

```
static struct super_operations ospfs_superblock_ops = {
};
```

// Functions used to hook the module into the kernel!

```
static int __init init_ospfs_fs(void)
{
    eprintk("Loading ospfs module...\n");
    return register_filesystem(&ospfs_fs_type);
}
```

```
static void __exit exit_ospfs_fs(void)
{
    unregister_filesystem(&ospfs_fs_type);
    eprintk("Unloading ospfs module\n");
}
```

```
module_init(init_ospfs_fs)
module_exit(exit_ospfs_fs)
```

```
// Information about the module
MODULE_AUTHOR("Jonathan Woong");
MODULE_DESCRIPTION("OSPFS");
MODULE_LICENSE("GPL");
```