

# WEENSY OS

```
#include "mpos-kern.h"
#include "x86.h"
#include "lib.h"

/*****
 * mpos-kern
 *
 * This is the miniprocs's kernel.
 * It sets up a process descriptor for the initial application, then runs
 * that application, and responds to its system calls.
 *
 *****/

// The kernel is loaded starting at 0x100000.
// The miniprocs applications are also available in RAM in packed form.
// The kernel loads one of those applications into memory starting at 0x200000.
// It also allocates 1/4 MB for each possible miniprocs's stack, starting at
// 0x280000.
// Each process's stack grows down from the top of its stack space.

#define PROC1_STACK_ADDR      0x280000
#define PROC_STACK_SIZE      0x040000

// MINIPROCOS MEMORY MAP
//
// +-----+-----+-----+-----+-----+
// | Base Memory (640K) | I/O Memory | Kernel | Kernel |
// | (unused) | | Code + Data | Stack |
// +-----+-----+-----+-----+
// | 0 | 0xA0000 | 0x100000 | 0x200000 |
// |
// | /+-----+-----+-----+-----+ /
// | | Application | Miniproc 1 | Miniproc 2 | Miniproc 3 |
// | | Code + Globals | Stack | Stack | Stack |
// | | /+-----+-----+-----+-----+ /
// | | 0x200000 | 0x280000 | 0x2C0000 | 0x300000 | 0x340000 |
// | | | | | | | |
// | | PROC1_STACK_ADDR | PROC1_STACK_ADDR |
// | | | + 2*PROC_STACK_SIZE |
// | | |
// | | PROC1_STACK_ADDR |
// | | + PROC_STACK_SIZE |
// |
// There is also a shared 'cursorpos' variable, located at 0x60000 in the
// kernel's data area. (This is used by 'app_printf' in mpos-app.h.)

// A process descriptor for each possible miniprocs.
// Note that proc_array[0] is never used.
// The main application process descriptor is proc_array[1].
static process_t proc_array[NPROCS];

// A pointer to the currently running process.
// This is kept up to date by the run() function, in mpos-x86.c.
process_t *current;

/*****
 * start
 *
 * Initialize the hardware and process descriptors to empty, then load
 * and run the first process.
 *
 *****/

void
start(void)
```

```
{
const char *s;
int whichprocess;
pid_t i;

// Initialize process descriptors as empty
memset(proc_array, 0, sizeof(proc_array));
for (i = 0; i < NPROCS; i++) {
proc_array[i].p_pid = i;
proc_array[i].p_state = P_EMPTY;
}

// The first process has process ID 1.
current = &proc_array[1];

// Set up x86 hardware, and initialize the first process's
// special registers. This only needs to be done once, at boot time.
// All other processes' special registers can be copied from the
// first process.
segments_init();
special_registers_init(current);

// Erase the console, and initialize the cursor-position shared
// variable to point to its upper left.
console_clear();

// Figure out which program to run.
cursorpos = console_printf(cursorpos, 0x0700, "Type '1' to run mpos-app, or '2' to run mpos-app2.");
do {
whichprocess = console_read_digit();
} while (whichprocess != 1 && whichprocess != 2);
console_clear();

// Load the process application code and data into memory.
// Store its entry point into the first process's EIP
// (instruction pointer).
program_loader(whichprocess - 1, &current->p_registers.reg_eip);

// Set the main process's stack pointer, ESP.
current->p_registers.reg_esp = PROC1_STACK_ADDR + PROC_STACK_SIZE;

// Mark the process as runnable!
current->p_state = P_RUNNABLE;

// Switch to the main process using run().
run(current);
}

/*****
 * interrupt
 *
 * This is the weensy interrupt and system call handler.
 * New system calls are implemented by code in this function.
 *
 *****/

static pid_t do_fork(process_t *parent);

void
interrupt(registers_t *reg)
{
// The processor responds to a system call interrupt by saving some of
// the application's state on the kernel's stack, then jumping to
// kernel assembly code (in mpos-int.S, for your information).
// That code saves more registers on the kernel's stack, then calls
// interrupt(). The first thing we must do, then, is copy the saved
```

## WEENSY OS

```
// registers into the 'current' process descriptor.
current->p_registers = *reg;
```

```
switch (reg->reg_intno) {
```

```
case INT_SYS_GETPID:
// The 'sys_getpid' system call returns the current
// process's process ID. System calls return results to user
// code by putting those results in a register. Like Linux,
// we use %eax for system call return values. The code is
// surprisingly simple:
current->p_registers.reg_eax = current->p_pid;
run(current);
```

```
case INT_SYS_FORK:
// The 'sys_fork' system call should create a new process.
// You will have to complete the do_fork() function!
current->p_registers.reg_eax = do_fork(current);
run(current);
```

```
case INT_SYS_YIELD:
// The 'sys_yield' system call asks the kernel to schedule a
// different process. (MiniprocOS is cooperatively
// scheduled, so we need a special system call to do this.)
// The schedule() function picks another process and runs it.
schedule();
```

```
case INT_SYS_EXIT:
// 'sys_exit' exits the current process, which is marked as
// non-runnable.
// The process stored its exit status in the %eax register
// before calling the system call. The %eax REGISTER has
// changed by now, but we can read the APPLICATION's setting
// for this register out of 'current->p_registers'.
```

```
current->p_state = P_EMPTY;
current->p_exit_status = current->p_registers.reg_eax;
```

```
if(current->p_wait_pid)
{
proc_array[current->p_wait_pid].p_registers.reg_eax = current->p_exit_status;
proc_array[current->p_wait_pid].p_state = P_RUNNABLE;
current->p_wait_pid = 0;
}
schedule();
```

```
case INT_SYS_WAIT: {
// 'sys_wait' is called to retrieve a process's exit status.
// It's an error to call sys_wait for:
// * A process ID that's out of range (<= 0 or >= NPROCS).
// * The current process.
// * A process that doesn't exist (p_state == P_EMPTY).
// (In the Unix operating system, only process P's parent
// can call sys_wait(P). In MiniprocOS, we allow ANY
// process to call sys_wait(P).)
```

```
pid_t p = current->p_registers.reg_eax;
if (p <= 0 || p >= NPROCS || p == current->p_pid
|| proc_array[p].p_state == P_EMPTY)
current->p_registers.reg_eax = -1;
else if (proc_array[p].p_state == P_ZOMBIE)
current->p_registers.reg_eax = proc_array[p].p_exit_status;
else{
current->p_state = P_BLOCKED;
proc_array[p].p_wait_pid = current->p_pid;
}
```

```
schedule();
```

```
}
```

```
default:
while (1)
/* do nothing */;
```

```
}
}
```

```
/******
* do_fork
*
* This function actually creates a new process by copying the current
* process's state. In MiniprocOS, a process's state consists of
* (1) its registers, and (2) its stack -- that's it. All processes share
* THE SAME code and global variables. (So really we should call them
* "miniprocesses" or something.)
* The parent process is passed in as an argument. The function should
* return the process ID of the child process, or -1 if the function can't
* create a child process.
* Your job is to fill it in!
*
*****/
```

```
static void copy_stack(process_t *dest, process_t *src);
```

```
static pid_t
do_fork(process_t *parent)
{
// YOUR CODE HERE!
// First, find an empty process descriptor. If there is no empty
// process descriptor, return -1. Remember not to use proc_array[0].
// Then, initialize that process descriptor as a running process
// by copying the parent process's registers and stack into the
// child. Copying the registers is simple: they are stored in the
// process descriptor in the 'p_registers' field. Copying the stack
// is a little more involved; see the copy_stack function, below.
// The child process's registers will be equal to the parent's, with
// two differences:
// * reg_esp The child process's stack pointer will point into
// its stack, rather than the parent's. copy_stack
// should arrange this.
// * ??????? There is one other difference. What is it? (Hint:
// What should sys_fork() return to the child process?)
// You need to set one other process descriptor field as well.
// Finally, return the child's process ID to the parent.
```

```
///// USER DEFINED VALUES
process_t *child = NULL;
int foundEmptyProcessDescriptor=0; // 0 = NO, 1 = YES
int i=1; // do not start at proc_array[0]
```

```
///// FIND EMPTY PROCESS DESCRIPTOR
for(i; i< NPROCS; i++) // loop through all process descriptors
{
if (proc_array[i].p_state == P_EMPTY) // if a process descriptor is empty
{
child = &proc_array[i]; // child points to that empty process descriptor
foundEmptyProcessDescriptor=1; // set flag
break;
}
}
```

```
///// FORK CHILD
if (foundEmptyProcessDescriptor)
{
```

## WEENSY OS

```
child->p_registers = parent->p_registers; // copy registers
copy_stack(child,parent); // copy stack
child->p_state = P_RUNNABLE; // set child as runnable
child->p_registers.reg_eax = 0; // return 0 to child
return child->p_pid; // parent receives child pid
}
else
{
return -1;
}
}
```

```
static void
copy_stack(process_t *dest, process_t *src)
{
uint32_t src_stack_bottom, src_stack_top;
uint32_t dest_stack_bottom, dest_stack_top;
```

```
// YOUR CODE HERE!
// This function copies the 'src' process's stack into the 'dest'
// process's stack region. Then it sets 'dest's stack pointer to
// correspond to 'src's stack pointer.
```

```
// For example, assume that 'src->p_pid == 1' and 'dest->p_pid == 2'.
// Then this code should change this memory setup:
```

```
//      Miniproc 1 Stack   Miniproc 2 Stack
//  /--+-+-----+-----+---/
//  |   ABXLQPAOSRJ|       |
//  /--+-+-----+-----+---/
//  | 0x280000   ^   0x2C0000   0x300000
//  |
//  | src->p_registers.reg_esp
//  | == 0x29A4CC
```

// into this:

```
//      Miniproc 1 Stack   Miniproc 2 Stack
//  /--+-+-----+-----+---/
//  |   ABXLQPAOSRJ|   ABXLQPAOSRJ|
//  /--+-+-----+-----+---/
//  | 0x280000   ^   0x2C0000   ^   0x300000
//  |           |           |
//  | src->p_registers.reg_esp   dest->p_registers.reg_esp
//  | == 0x29A4CC           == 0x2DA4CC
//  |           == 0x300000 + (0x29A4CC - 0x2C0000)
```

// You may implement this however you like, but we found it easiest  
// to express with variables that locate the bottom and top of each  
// stack. In our examples, the variables would equal:

```
//  /--+-+-----+-----+---/
//  |   ABXLQPAOSRJ|   ABXLQPAOSRJ|
//  /--+-+-----+-----+---/
//  | 0x280000   ^   0x2C0000   ^   0x300000
//  |           |           |           ^
//  | src_stack_bottom | dest_stack_bottom |
//  | == 0x29A4CC   | == 0x2DA4CC   |
//  |           src_stack_top   dest_stack_top
//  |           == 0x2C0000   == 0x300000
```

// Your job is to figure out how to calculate these variables,  
// and then how to actually copy the stack. (Hint: use memcpy.)  
// We have done one for you.

```
////// CALCULATE STACK VARIABLES
src_stack_top = PROC1_STACK_ADDR + (PROC_STACK_SIZE) * (src->p_pid);
src_stack_bottom = src->p_registers.reg_esp;
```

```
dest_stack_top = PROC1_STACK_ADDR + (dest->p_pid)*PROC_STACK_SIZE;
dest_stack_bottom = dest_stack_top + (src_stack_bottom - src_stack_top);
```

```
////// COPY STACK
memcpy((void *)dest_stack_bottom,(void *)src_stack_bottom, (src_stack_top-src_stack_bottom));
////// SET DEST ESP
dest->p_registers.reg_esp = dest_stack_bottom;
}
```

```
/******
* schedule
*
* This is the weensy process scheduler.
* It picks a runnable process, then context-switches to that process.
* If there are no runnable processes, it spins forever.
*
*****/
```

```
void
schedule(void)
{
pid_t pid = current->p_pid;
while (1) {
pid = (pid + 1) % NPROCS;
if (proc_array[pid].p_state == P_RUNNABLE)
run(&proc_array[pid]);
}
}
```

## WEENSY OS

```
#ifndef WEENSYOS_MPOS_KERN_H
#define WEENSYOS_MPOS_KERN_H
#include "mpos.h"
#include "x86.h"

// Process state type
typedef enum procstate {
    P_EMPTY = 0,                // The process table entry is empty
    // (i.e. this is not a process)
    P_RUNNABLE,                // This process is runnable
    P_BLOCKED,                // This process is blocked
    P_ZOMBIE                    // This process has exited, but no one
    // has called sys_wait() yet
} procstate_t;

// Process descriptor type
typedef struct process {
    pid_t p_pid;                // Process ID

    registers_t p_registers;    // Current process state: registers,
    // stack location, EIP, etc.
    // 'registers_t' defined in x86.h
    procstate_t p_state;        // Process state; see above
    int p_exit_status;          // Process's exit status (if it has
    // exited and p_state == P_ZOMBIE)

    pid_t p_wait_pid; // pid that we wait on
} process_t;

// Top of the kernel stack
#define KERNEL_STACK_TOP      0x80000

// Functions defined in mpos-kern.c
void interrupt(registers_t *reg);
void schedule(void);

// Functions defined in mpos-x86.c
void segments_init();
void special_registers_init(process_t *proc);
void console_clear(void);
int console_read_digit(void);
// Function defined in mpos-loader.c
void program_loader(int programnumber, uint32_t *entry_point);

extern process_t *current;
void run(process_t *proc) __attribute__((noreturn));

#endif
```